


1. Resize image.
2. Run convolutional network.
3. Non-max suppression.




Object Detection - YOLO

Literature Sharing: Xiaokang Guo

Term Project

- Get SMILES


- C-C(=O)-O-C(-C-C(=O)-O)-C-N(-C)(-C)-C
- Presence of Atom → Character
- Atom Type → C, N, O, S, ...
- Presence of Bond → Symbol
- Bond Type → = or #





Term Project

Presence of Atom
Atom Type
Presence of Bond
Bond Type

Unbonded





Bonded



Sliding
Window

Capability of Object Detector

- Add bounding box and class identifier simultaneously




You Only Look Once Version 1

- Fast end-to-end CNN
 - Fast enough for real-time detection
 - One shot object detector
- Other mentioned models ?
 - Region CNN
 - Fast R-CNN
 - Faster R-CNN


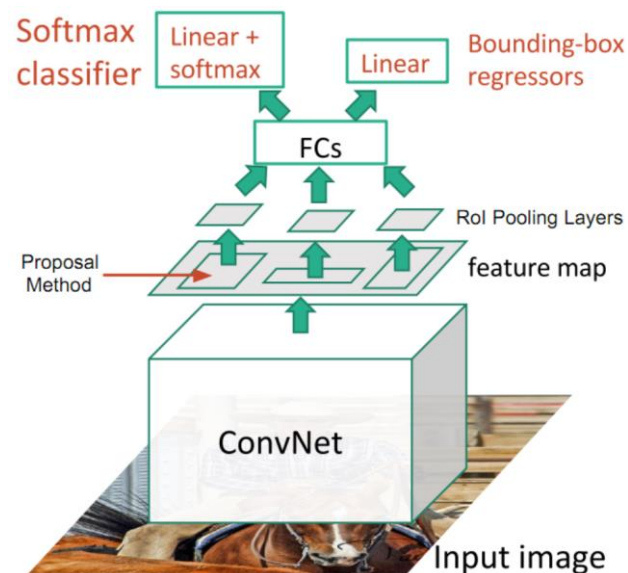
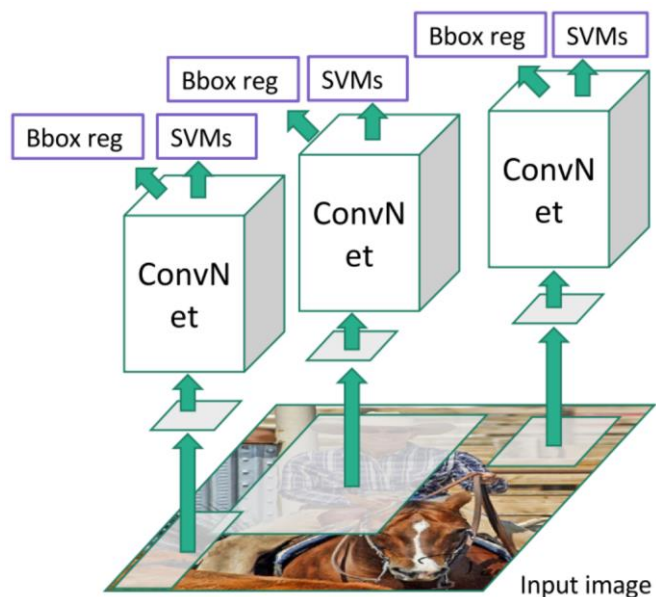
Region Proposal

- Selective Search (SS):
 1. Generate initial sub-segmentation, we generate many candidate regions
 2. Use greedy algorithm to recursively combine similar regions into larger ones
 3. Use the generated regions to produce the final candidate region proposals




R-CNN / Fast R-CNN / Faster R-CNN


- Image \rightarrow SS \rightarrow Region Proposal (2k) \rightarrow CNN \rightarrow Features \rightarrow SVM \rightarrow Objectness
 - 47s / image with VGG16 [Simonyan & Zisserman. ICLR15]
- Image \rightarrow ConvL \rightarrow Features \rightarrow SS \rightarrow Proposal \rightarrow RoI Pooling \rightarrow FCL \rightarrow Result
 - Around 1 second per image
- Proposal Method \Rightarrow Region Proposal Network




Region Proposal Network



8 x 8 feature maps



3 proposals for each location



You Only Look Once

- Arbitrary Size



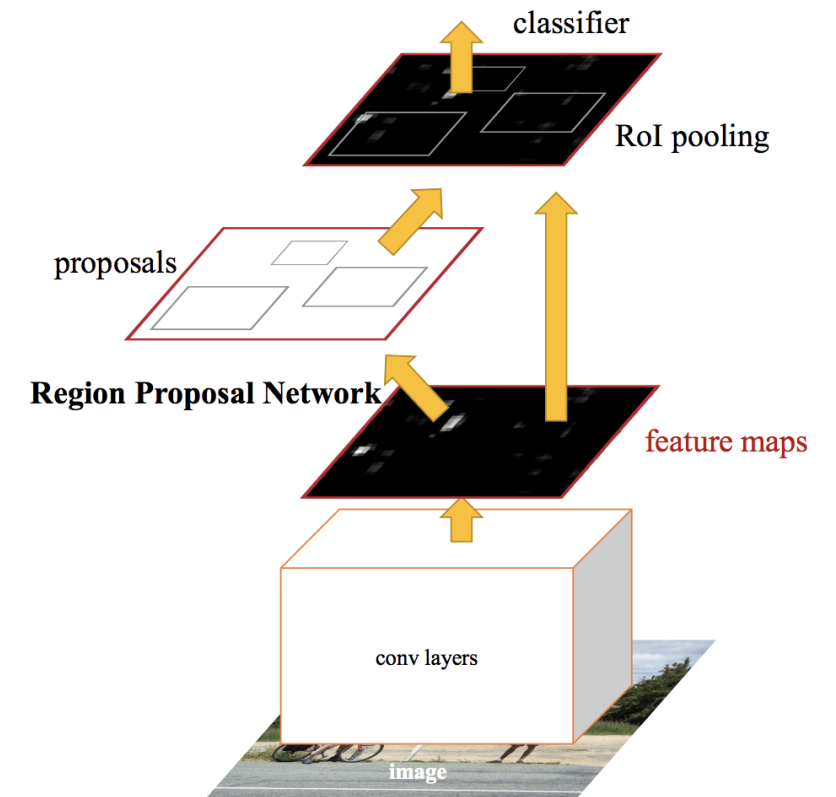

You Only Look Once


- Fixed Size (448px×448px)



Capability

- Localized detector which might be responsible for the whole image





	Conv 1	Max Pool 1	Conv 2	Max Pool 2
Filters	(7,7,64,2)	(2,2,2)	(3,3,192)	(2,2,2)
Output	224×224×64	112×112×64	112×112×192	56×56×192


Conv 3	Conv 4	Conv 5	Conv 6	Max Pool 3
(1,1,128)	(3,3,256)	(1,1,256)	(1,1,512)	(2,2,2)
56×56×128	56×56×256	56×56×256	56×56×512	28×28×512

	C7	C8	C9	C10	C11	C12	C13	C14	C15	C16	P4
Filters	(1,1,256)	(3,3,512)	(1,1,256)	(3,3,512)	(1,1,256)	(3,3,512)	(1,1,256)	(3,3,512)	(1,1,512)	(3,3,1024)	(2,2,2)
Output	28×28×256	28×28×512	28×28×256	28×28×512	28×28×256	28×28×512	28×28×256	28×28×512	28×28×512	28×28×1024	14×14×1024


	C17	C18	C19	C20	C21	C22	C23	C24	F1
Filters	(1,1,512)	(3,3,1024)	(1,1,512)	(3,3,1024)	(3,3,1024)	(3,3,1024,2)	(3,3,1024)	(3,3,1024)	-
Output	14×14×512	14×14×1024	14×14×512	14×14×1024	14×14×1024	7×7×1024	7×7×1024	7×7×1024	4096

F2
-
7×7×30

Output Layer



- Grid (7×7), Cells (3,1), Bounding boxes (2), Prediction Vector (30)



One predictions vector for each cell (49 cells in total):

$x_1, y_1, w_1, h_1, c_1, x_2, y_2, w_2, h_2, c_2, C_1, C_2, C_3, \dots, C_{20}$ PASCAL VOC

confidence score $c_i = \text{Pr}(\text{Obj}) * \text{IOU}_{\text{pred}}^{\text{truth}}$ Objectness and accuracy

conditional class probability $C_n = \text{Pr}(\text{Class}_n | \text{Obj}) = p_i(n)$

The prob of belonging to n -th class if an obj is in i -th cell.

First Bounding Box

$$x_1 = (250 - 189) / 63 = 0.968$$

$$y_1 = (85 - 63) / 63 = 0.349$$

$$w_1 = (124 - 60) / 448 = 0.143$$

$$h_1 = (378 - 120) / 448 = 0.576$$

At test time, for each box the confidence score times with conditional class probability gives class-specific confidence scores: $c_i * C_n = \text{Pr}(\text{Class}_n | \text{Obj}) * \text{Pr}(\text{Obj}) * \text{IOU}_{\text{pred}}^{\text{truth}}$

Loss Function

$x_1, y_1, w_1, h_1, c_1, x_2, y_2, w_2, h_2, c_2, C_1, C_2, C_3, \dots, C_{20}$

$$\begin{aligned}
 \mathbf{RSS} = & \lambda_{coord} \sum_{i=0}^{S^2} \sum_{j=0}^B \mathbb{1}_{ij}^{obj} [(x_i - \hat{x}_i)^2 + (y_i - \hat{y}_i)^2] + \lambda_{coord} \sum_{i=0}^{S^2} \sum_{j=0}^B \mathbb{1}_{ij}^{obj} [(\sqrt{w_i} - \sqrt{\hat{w}_i})^2 + (\sqrt{h_i} - \sqrt{\hat{h}_i})^2] \\
 & + \sum_{i=0}^{S^2} \sum_{j=0}^B \mathbb{1}_{ij}^{obj} (c_i - \hat{c}_i)^2 + \lambda_{no_obj} \sum_{i=0}^{S^2} \sum_{j=0}^B \mathbb{1}_{ij}^{no_obj} (c_i - \hat{c}_i)^2 + \sum_{i=0}^{S^2} \mathbb{1}_i^{obj} \sum_{c \in classes} (p_i(c) - \hat{p}_i(c))^2
 \end{aligned}$$


$$\lambda_{coord} = 5, \quad \lambda_{no_obj} = 0.5, \quad \mathbb{1}_{ij}^{obj} = \begin{cases} 1 & , \text{obj in} \\ 0 & , \text{obj out} \end{cases}, \quad \mathbb{1}_{ij}^{no_obj} = \begin{cases} 0 & , \text{obj in} \\ 1 & , \text{obj out} \end{cases}$$

References





- YOLO v1 Original Paper
 - <https://arxiv.org/abs/1506.02640>
- Understanding YOLO
 - <https://hackernoon.com/understanding-yolo-f5a74bbc7967>
- Selective Search
 - <https://www.koen.me/research/pub/uijlings-ijcv2013-draft.pdf>
- R-CNN, Fast R-CNN, Faster R-CNN, YOLO – Object Detection Algorithm
 - <https://towardsdatascience.com/r-cnn-fast-r-cnn-faster-r-cnn-yolo-object-detection-algorithms-36d53571365e>
- RPN
 - https://medium.com/@jonathan_hui/what-do-we-learn-from-region-based-object-detectors-faster-r-cnn-r-fcn-fpn-7e354377a7c9
- F.F. Li et., CS231n, Lecture 11, Stanford, 2017
 - http://cs231n.stanford.edu/slides/2017/cs231n_2017_lecture11.pdf

Thanks for your attention!


RoI Pooling in Fast R-CNN



Region of Interest Pooling Layer

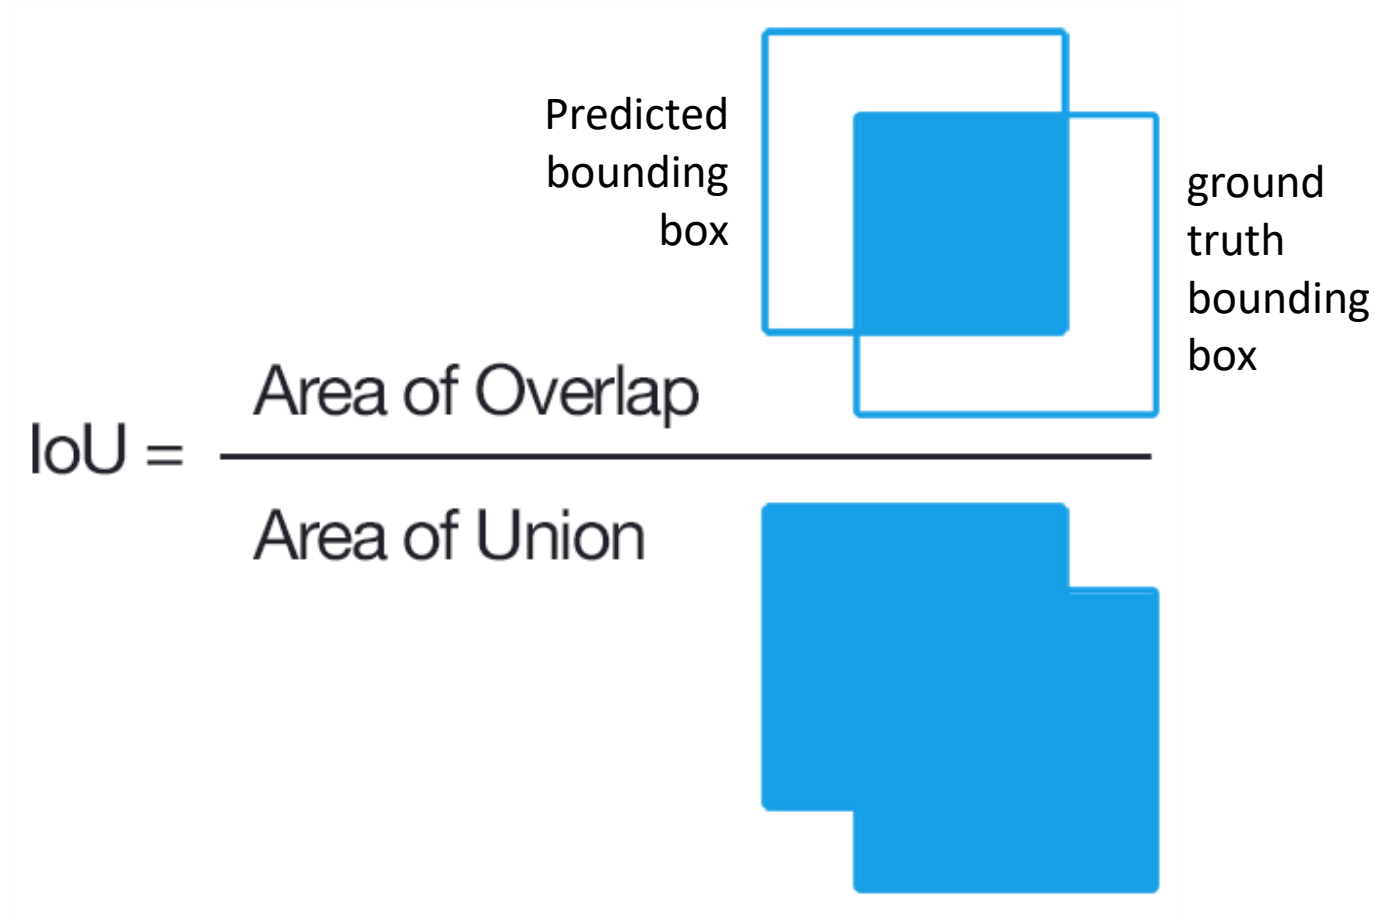


BBox Regressor




Use regression to refine the original ROI in blue to the red one

Intersection over union



1×1 Convolutional Layer



Andrew Ng

<https://www.youtube.com/watch?v=vcp0XvDAX68>

Train Strategy

- The authors describe the training in the following way
 - First, pretrain the first 20 convolutional layers using the ImageNet 1000-class competition dataset, using a input size of 224x224
 - Then, increase the input resolution to 448x448
 - Train the full network for about 135 epochs using a batch size of 64, momentum of 0.9 and decay of 0.0005
 - Learning rate schedule: for the first epochs, the learning rate was slowly raised from 0.001 to 0.01. Train for about 75 epochs and then start decreasing it.
 - Use data augmentation with random scaling and translations, and randomly adjusting exposure and saturation.